

Pushing Data through Pipelines with Coroutines



Axel Sirota

MACHINE LEARNING ENGINEER

@AxelSirota

Much More Than Just Iteration: Generator Based Data Pipeline

A Pipeline Example

```
In [9]: def pipeline(number):
...:     data = (i for i in range(number))
...:     squared = (i**2 for i in data)
...:     negated = (-i for i in squared)
...:     return (n + 1 for n in negated)
...:

In [10]: list(pipeline(10))
Out[10]: [1, 0, -3, -8, -15, -24, -35, -48, -63, -80]
```

- ◀ # This is **NOT** a generator, but a **generator factory**
- ◀ # Get a generator that yields ints; this is **range**
- ◀ # Square each element of data *when needed*
- ◀ # Negate each element in squared.
- ◀ # Add 1 to each element. This is $-X^2 + 1$

Note that until we iterate over this generator, we don't execute!

A Pipeline Example

```
In [15]: def squared(iterable):  
...:     return (i**2 for i in iterable)  
...:
```

```
In [16]: def negated(iterable):  
...:     return (-i for i in iterable)  
...:
```

```
In [17]: def add_one(iterable):  
...:     return (i + 1 for i in iterable)  
...:
```

```
In [18]: def pipeline(number):  
...:     return  
add_one(negated(squared(range(number))))  
...:
```

```
In [19]: list(pipeline(10))  
Out[19]: [1, 0, -3, -8, -15, -24, -35, -48, -63, -80]
```


◀ **# This makes sense on its own!**

◀ **# Same here! We just negate an iterable (Given it supports `__add__`)**


◀ **# Same here!**

◀ **# We just need to chain the generators together to return another generator**


A Pipeline Example




Generator Based Pipeline: A Review



0 Storage -> it scales to infinity



Is like unix pipelines, we can chain generators together



To learn more -> <http://www.dabeaz.com/generators/Generators.pdf>



Generators can be chained to create pipelines of execution

They decouple the definition of a task from its execution

`('THE' + ' ' + 'PROGRAMMERS') ...`

Why I Don't Have to Remember All This: Itertools!

A Test For Itertools

```
In [2]: sample = [5,4,2,8,7,6,3,0,9,1]
```

```
In [3]: list(averager(sample))
```

```
Out[3]:
```

```
[5.0,  
4.5,  
3.6666666666666665,  
4.75,  
5.2,  
5.333333333333333,  
5.0,  
4.375,  
4.888888888888889,  
4.5]
```

◀ **# Given a sample ...**

◀ **# We print the average up to a point**

◀ **# The average of 5 is 5**

◀ **# The average of 5 and 4 is 4.5**

◀ **The average of 5,4,2 and 8 is 4.75 and
so on...**

The Old Way

```
def old_style_averager(iterable):  
    total_sum = 0  
    total_elements = 0  
    average = []  
    for number in iterable:  
        total_sum += number  
        total_elements += 1  
        average.append(total_sum/  
total_elements)  
    return average
```

◀ # Given a sample ...

◀ # We iterate over it

◀ # Keep the updated average

The Old Way

```
In [10]: sample =  
np.random.random(10000)
```

```
In [11]: %timeit  
old_style_averager(sample)  
3.33 ms ± 63 µs per loop (mean ± std. dev.  
of 7 runs, 100 loops each)
```

```
In [12]: %load_ext memory_profiler
```

```
In [13]: %memit  
old_style_averager(sample)  
peak memory: 71.22 MiB, increment: 0.20  
MiB
```

◀ **# Not bad!**

◀ **# Keep a cap on RAM too**

The Generator Way

```
def generator_averager(iterable):  
    total_sum = 0  
    total_elements = 0  
    for number in iterable:  
        total_sum += number  
        total_elements += 1  
    yield total_sum/total_elements
```

◀ # Given a sample ...

◀ # We iterate over it

◀ # But we yield at each point!

If we need a middle average, we don't iterate over all

The Generator Way

```
In [27]: sample =  
np.random.random(10000)
```

```
In [28]: %timeit  
generator_averager(sample)  
214 ns ± 2.31 ns per loop (mean ± std. dev.  
of 7 runs, 1000000 loops each)
```

```
In [29]: %memit  
generator_averager(sample)  
peak memory: 73.80 MiB, increment:  
-0.08 MiB
```

◀ **# 94% Faster!**

◀ **# Same memory usage!**

Itertools Magic

```
In [24]: for i in itertools.starmap(pow,  
...:    [(2,5), (3,2), (10,3)]):  
...:     print(i)  
...:  
32  
9  
1000
```

- **Applies the function to every tuple in the iterable!**
- **So $1000 = \text{pow}(10,3)$**
- **We could apply the division if we could get a tuple that accumulates the sum and number of elements...**

Itertools Magic

```
In [26]: for i in
itertools.accumulate([2,3,4,5,6], lambda
a,b: a*b):
...:     print(f'-->{i}')
...:
...:
-->2
-->6
-->24
-->120
-->720
```

- **Applies the function to the first element, pairing with 1**
- **Then it applies the result of that (2) to the next (3)**
- **Then it applies the result of that (6) to the next element (4)**
- **We could get the sums and total number of elements if we apply to tuples...**

Itertools Magic

```
In [28]: for i in
enumerate(itertools.accumulate([2,3,4,5,6
]), 1):
    ...:     print(f'-->{i}')
    ...:
    ...:
-->(1, 2)
-->(2, 5)
-->(3, 9)
-->(4, 14)
-->(5, 20)
```

- **Enumerate counts the number of elements by 1**
- **So each of these tuples, if we divide them, are the average up to that point!**

Itertools Magic

```
In [29]: for i in itertools.starmap(lambda
a,b: b/a,
enumerate(itertools.accumulate([2,3,4,5,6
]), 1)):
....:     print(f'-->{i}')
....:
....:
....:
-->2.0
-->2.5
-->3.0
-->3.5
-->4.0
```

```
def averager(iterable):
    return itertools.starmap(lambda a, b: b / a,
enumerate(itertools.accumulate(iterable), 1))
```

- **Combining everything we get the expected result**
- **We got our average as 1 liner!**

Itertools Magic

```
In [27]: sample =  
np.random.random(10000)
```

```
In [28]: %timeit averager(sample)  
189 ns ± 1.12 ns per loop (mean ± std. dev.  
of 7 runs, 1000000 loops each)
```

```
In [29]: %memit averager(sample)  
peak memory: 43.12 MiB, increment: -0.08  
MiB
```

◀ **# 12% Faster than generator solution**

◀ **# Half the memory usage!**



Itertools offer generator factories

We can chain them to get pretty complex generators as one liners

They are optimized in C, so they are:

- Highly scalable**
- Minimal memory requirements**
- Recommended all times!**

Demo

Create a generator based pipeline to parse logs

Migrate much as possible to itertools to learn more on the module

Summary

Generators can be chained as pipeline to act deferred on a whole dataset

Itertools offer a lot of precooked generator factories

They can chained as a pipeline to generate pretty amazing results